



# Metadata Analysis with Zeek



# Metadata Analysis with Zeek

1. Zeek Overview
2. Installation & Configuration
3. Displaying and Filtering Logs

# Metadata Analysis with Zeek

1. Zeek Overview
2. Installation & Configuration
3. Displaying and Filtering Logs


# What is Zeek?

- Passive, open-source traffic analysis platform
- Most often used in cybersecurity, but helpful for performance monitoring and troubleshooting
- Generates rich set of logs that summarize each connection
- Port-independent protocol analysis
- Interactive tutorial available at <https://try.zeek.org/>

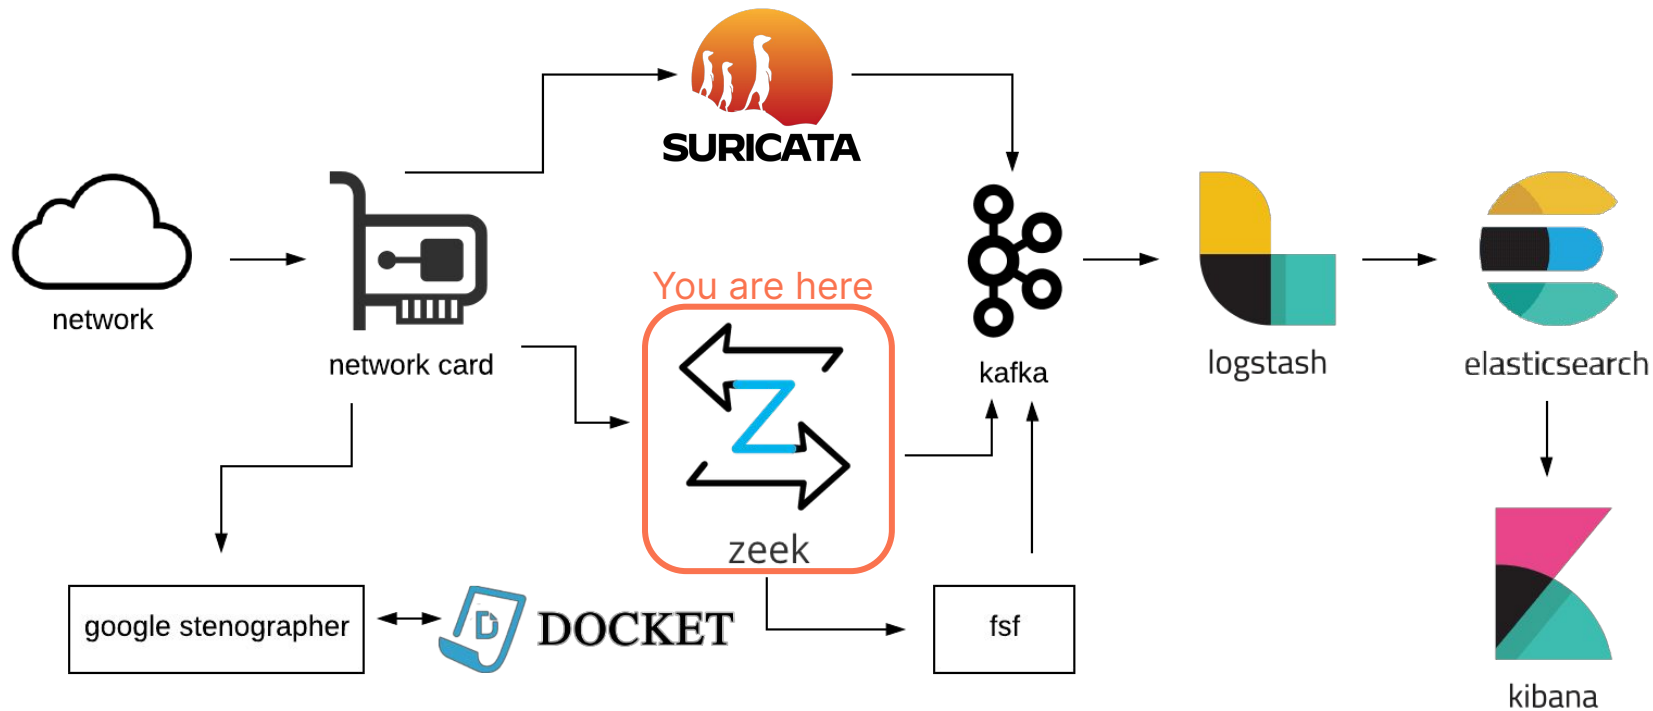


# History

- Created in 1995 by Vern Paxson at Lawrence Berkeley National Laboratory (LBNL) out of research interests
- Previously known as “Bro”, renamed to “Zeek” in 2018
- Heavily used in energy and super-computer research communities since late 90’s



# RockNSM



# Typical Network log data



---

Packet captures  
(PCAP)



**SURICATA**

---

Alerts



---

Session and Protocol  
Metadata Logs

# Packet Captures (PCAP)

- Pros
  - Allows complete reconstruction of network events
  - Contains all the raw data
- Cons
  - Tremendous amount of storage required
  - Does not scale well



# Alerts

- Pros
  - Good for finding “known bad things”
  - Automated alert generation given a rule has been matched
- Cons
  - Prone to false positives and negatives
  - Does not detect events you do not have rules for

# Zeek Logs - Session & Protocol Metadata

- Analyzes network data and creates session logs
- Can be used to construct full timeline of events
- See the bigger picture, especially retroactively!
- Port independent protocol analysis
- Uses the terms Originator and Responder
  - originator  $\neq$  source
  - responder  $\neq$  destination

# Packet vs Network Events

## Packet Events (Source/Destination)

10.0.0.1 (S)	→	200.1.1.6 (D)
10.0.0.1 (D)	←	200.1.1.6 (S)
10.0.0.1 (S)	→	200.1.1.6 (D)
10.0.0.1 (D)	←	200.1.1.6 (S)

- Based on individual packets
- Source and destination change based on the **direction** of the **communication**

TCPdump  
Wireshark  
Suricata

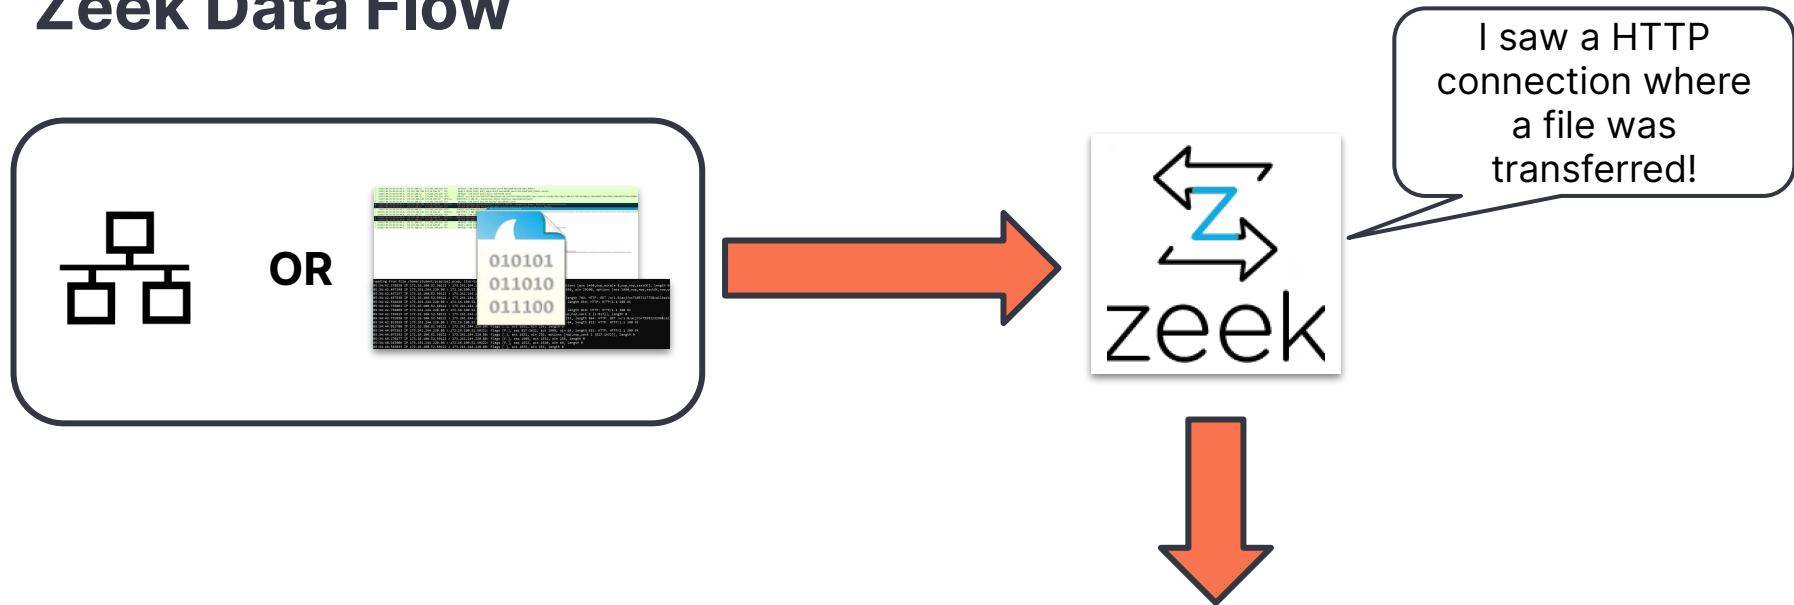
## Network Events (Client/Server)

Client (Originator)		Server (Responder)
10.0.0.1	→	200.1.1.6
10.0.0.1	←	200.1.1.6
10.0.0.1	→	200.1.1.6
10.0.0.1	←	200.1.1.6

- Based on who **initiated** the conversation
- Client == Originator
- Server == Responder
- Client and Server **do not change** over the course of the connection

Zeek

# Zeek Data Flow



conn.log

```
conn.log:1493012082.378828 CpTzSV12rEiY8zSgHk 172.16.100.52 50122 173.241.244.220 80 tcp http
5.906178 1608 1631 SF T F 0 ShAdDtFf 10 2044 6 3514 - - -
00:50:56:98:40:e6 00:50:56:98:6e:6c
```

http.log

```
http.log:1493012082.457538 CpTzSV12rEiY8zSgHk 172.16.100.52 50122 173.241.244.220 80 1 GET reuters-d.openx.net /w/1.0/acj?o=7105712775&callback
=OX_7105712775&ju=http://www.reuters.com/8jw-8b-38rf=0&aid=53 http://www.reuters.com/ 1.1 Mozilla/5.0 (Windows NT 6.1; Win64; x64) AppleWebKit/537.36 (KHTML, like G
ecko) Chrome/57.0.2987.133 Safari/537.36 -0 491 200 OK - - (empty) - -
FWYDvv18wXoaGbefhf -text/plain
```

files.log

```
files.log:1493012083.812839 F4nJFyleeX4UkCuxqd 173.241.244.220 172.16.100.52 CpTzSV12rEiY8zSgHk HTTP 0SHA1_MD5 text/plain - 0.000000
F F 491 - 0 0 F - 755554fff00b89c583bc4a2fb9722bca 8ac29ffbdb051f8dd7e761501a1500153ef1de68 - - -
```

# Zeek Logs

- **Network logs**
- File logs
- netControl
- Detection
- Observations
- Miscellaneous
- Diagnostic

# Metadata Analysis with Zeek

1. Zeek Overview
2. Installation & Configuration
3. Displaying and Filtering Logs

# Zeek Installation

- Zeek can be installed from source, or by using pre-built binary release packages for Linux.
  - Before Zeek can be installed, the required dependencies need to be installed on the target system.
  - Zeek can also make use of optional dependencies, if they are found at build time.
- In our lab environment, we are using the Zeek RPM directly provided by the official Zeek website.
- Can also be run on Docker using the official Docker images.



## Zeek Configuration

Path	Description
/usr/bin	<ul style="list-style-type: none"><li>• Binaries</li><li>• The application that will run</li></ul>
/opt/zeek/etc	<ul style="list-style-type: none"><li>• Configurations</li><li>• All the settings that manage Zeek preferences</li></ul>
/opt/zeek/logs	<ul style="list-style-type: none"><li>• Data</li><li>• Default directory where the data logs are stored</li></ul>
/opt/zeek/share/zeek/site	<ul style="list-style-type: none"><li>• Scripts</li><li>• Custom scripts allowing for customization</li></ul>



# /opt/zeek/etc

## node.cfg

- Default - standalone
- Used for Zeek clusters

## networks.cfg

- Defines the local networks

## zeekctl.cfg

- Log Rollover
- Log Path
- Config Directory

# Zeek Log Structure

- Defined in the **LogDir** parameter in `/opt/zeek/etc/zeekctl.cfg`

## `/opt/zeek/logs/current`

```
conn.log  
dhcp.log  
dns.log  
http.log  
...
```

## `/opt/zeek/logs`

2022-04-15

2022-04-16

2022-04-17

2022-04-18

2022-04-19

`current` -> `/opt/zeek/spool`

Note: Connections are only logged when the connection closes. If a connection is open when the rollover period occurs, it will not be in that log. See [corelight/zeek-long-connections](https://corelight.com/docs/zeek-long-connections) for more information.

# Metadata Analysis with Zeek

1. Zeek Overview
2. Installation & Configuration
3. Displaying and Filtering Logs

# Zeek Logs

- Plain ASCII text files
- Column-ized data - `<tab>` delimited
- Module Intent:
  - Need to understand source data
  - Command line tools to filter
- For use with Elasticsearch, logs can be output in json format
  - Final destination:
    - Through Logstash
    - Into Elasticsearch
    - Viewed in Kibana

# Zeek CLI - Command Options

- Getting basic help commands

```
zeek --help
```

- Running Zeek
  - C = ignore checksums
  - r = read from a tcpdump or pcap file

```
zeek -Cr [pcap file]
```

# **Labs 1 & 2**

## **Zeek: Functions Check**

## **Zeek: Additional Configuration**

# RockNSM Scripts

# Rock Bash Aliases

- Available on GitHub in [rocknsm/rock](https://github.com/rocknsm/rock)
  - /etc/profile.d/rock.sh
  - Included in lab environment
- [lesscolor](#)
  - performs a less, and *colorizes* each column

```
lesscolor conn.log
```

- [fields](#)
  - Lists out each column header

```
fields conn.log
```



# lesscolor

#ts	uid	id.orig_h	id.orig_p	id.resp_h	id.resp_p	proto	service	duration	orig_bytes	resp_bytes	conn_state	local_orig	local_resp						
#time	string	addr	port	addr	port	enum	string	interval	count	count	string	bool	bool	count	string	count	count	count	
1258531221.486539	CI6NdH3ZB2v7xBYdJ2	192.168.1.102	68	192.168.1.1	67	udp	dhcp	0.163820	301	300	SF	-	-	0	Dd	1	329	1	328
1258531680.237254	COAPnR2Pm5Mfsd18S2	192.168.1.103	137	192.168.1.255	137	udp	dns	3.780125	350	0	S0	-	-	0	D	7	546	0	-
1258531693.816224	CJYCE81MGL6aYS3kQ1	192.168.1.102	137	192.168.1.255	137	udp	dns	3.748647	350	0	S0	-	-	0	D	7	546	0	-
1258531635.800933	CqCtZg4uzGGdgaQV79	192.168.1.103	138	192.168.1.255	138	udp	-	46.725380	560	0	S0	-	-	0	D	3	644	0	-
1258531693.825212	CKhSPdIU8qik9TgVD1	192.168.1.102	138	192.168.1.255	138	udp	-	2.248589	348	0	S0	-	-	0	D	2	404	0	-
1258531803.872834	CXapBL3tYPRY1h0uai	192.168.1.104	137	192.168.1.255	137	udp	dns	3.748893	350	0	S0	-	-	0	D	7	546	0	-
1258531747.077012	CLW3nW2FSoPqT1eLq6	192.168.1.104	138	192.168.1.255	138	udp	-	59.052898	549	0	S0	-	-	0	D	3	633	0	-
1258531924.321413	CHF5fsFlg3S5xrkvk	192.168.1.103	68	192.168.1.1	67	udp	dhcp	0.044779	303	300	SF	-	-	0	Dd	1	331	1	328
1258531939.613071	CYF0oF2h7f66TszH2k	192.168.1.102	138	192.168.1.255	138	udp	-	-	-	S0	-	-	0	D	1	229	0	-	
1258532046.693816	CF00hdldmNj5tyXjif	192.168.1.104	68	192.168.1.1	67	udp	dhcp	0.002103	311	300	SF	-	-	0	Dd	1	339	1	328
1258532143.457078	C58MAO2cemWA8TdzcF	192.168.1.102	1170	192.168.1.1	53	udp	dns	0.068511	36	215	SF	-	-	0	Dd	1	64	1	243
1258532203.657268	CHwIU34Z4r9ais4QP8	192.168.1.104	1174	192.168.1.1	53	udp	dns	0.170962	36	215	SF	-	-	0	Dd	1	64	1	243
1258532331.365294	C9jNLV3Drof2P6kbu2	192.168.1.1	5353	224.0.0.251	5353	udp	dns	0.100381	273	0	S0	-	-	0	D	2	329	0	-
1258532331.365330	CHoU8E1Hxj3eSqNou5	fe80::219:e3ff:fee7:5d23	5353	ff02::fb	5353	udp	dns	0.100371	273	0	S0	-	-	0	D	2	369	0	-
1258532404.734264	Ch7dgg2GLGcJgwiQN6	192.168.1.103	137	192.168.1.255	137	udp	dns	3.873818	350	0	S0	-	-	0	D	7	546	0	-
1258532418.272517	C4YN5o49GYASXEYAl4	192.168.1.102	137	192.168.1.255	137	udp	dns	3.748891	350	0	S0	-	-	0	D	7	546	0	-
1258532404.859431	CtVTPq21CKZUANAfFk	192.168.1.103	138	192.168.1.255	138	udp	-	2.257840	348	0	S0	-	-	0	D	2	404	0	-
1258532456.089023	C7LWYL3fgOUOX5QCUL	192.168.1.102	1173	192.168.1.1	53	udp	dns	0.000267	33	497	SF	-	-	0	Dd	1	61	1	525
1258532418.281002	CPkDOJ2acDdhkEkMtb	192.168.1.102	138	192.168.1.255	138	udp	-	2.248843	348	0	S0	-	-	0	D	2	404	0	-
1258532525.592455	Cscgl1n4MWKPkbfLWj	192.168.1.1	5353	224.0.0.251	5353	udp	dns	0.099824	273	0	S0	-	-	0	D	2	329	0	-
1258532525.592493	Cly52n4KYzhOQL89Zg	fe80::219:e3ff:fee7:5d23	5353	ff02::fb	5353	udp	dns	0.099813	273	0	S0	-	-	0	D	2	369	0	-
1258532528.348891	C832Wd2HWTBfcCHKHQA	192.168.1.104	137	192.168.1.255	137	udp	dns	3.748895	350	0	S0	-	-	0	D	7	546	0	-
1258532528.357385	CEFY9r49EpntBcO149	192.168.1.104	138	192.168.1.255	138	udp	-	2.248339	348	0	S0	-	-	0	D	2	404	0	-
1258532644.128655	CwaEyG3DnZbJphU7E7	192.168.1.1	5353	224.0.0.251	5353	udp	dns	-	-	-	S0	-	-	0	D	1	154	0	-
1258532644.128680	CRZMsR68qg1CjCAoe	fe80::219:e3ff:fee7:5d23	5353	ff02::fb	5353	udp	dns	-	-	-	S0	-	-	0	D	1	174	0	-
1258532657.288677	CNKKNsQGictnKw5Ue	192.168.1.102	138	192.168.1.255	138	udp	-	-	-	-	S0	-	-	0	D	1	229	0	-
1258532683.876479	C83I0n4dSg4ywO8N95	192.168.1.103	138	192.168.1.255	138	udp	-	-	-	-	S0	-	-	0	D	1	240	0	-
1258532824.338291	Ctm1EV61s6Jkjmj42	192.168.1.104	138	192.168.1.255	138	udp	-	-	-	-	S0	-	-	0	D	1	229	0	-
1258533003.551468	CV7oiJ12P4YDaU4OX3	192.168.1.102	68	192.168.1.1	67	udp	dhcp	0.011807	301	300	SF	-	-	0	Dd	1	329	1	328
1258533129.324984	CAGsw43C5j5zqtDjvl	192.168.1.103	137	192.168.1.255	137	udp	dns	3.748641	350	0	S0	-	-	0	D	7	546	0	-
1258533142.729062	CKMGtM3LddiB1ghlz6	192.168.1.102	137	192.168.1.255	137	udp	dns	3.748893	350	0	S0	-	-	0	D	7	546	0	-
1258533129.333980	C2GhUk3IjzhZg4Aqs5	192.168.1.103	138	192.168.1.255	138	udp	-	2.248336	348	0	S0	-	-	0	D	2	404	0	-
1258533142.737803	C0KQMm2wzu5p7dkY47	192.168.1.102	138	192.168.1.255	138	udp	-	2.248086	348	0	S0	-	-	0	D	2	404	0	-

# Columns == Fields

**\$ fields conn.log**

```
1  ts
2  uid
3  id.orig_h
4  id.orig_p
5  id.resp_h
6  id.resp_p
7  proto
8  service
...
```

# Linux Command Line Interface

# Linux Command Line Interface (CLI)

- Use commands of your choice to view logs!
  - cat
  - less
  - head
- Commands of your choice to manipulate logs!
  - cut
  - awk
  - sort
  - uniq
- Useful resource: <https://linuxjourney.com/>

# Awk - What is it?

- Scans a file line by line
- Splits each input line into fields
- Compares input line/fields to pattern
- Performs action(s) on matched lines

## Syntax:

```
awk options 'selection _criteria {action }' input-file > output-file
```

# 'cut' with 'awk'


```
#path conn
#open 2020-10-06-17-34-39
#ts uid id.orig_h id.orig_p id.resp_h id.resp_p proto service duration orig_bytes resp
#time string addr port addr port enum string interval count count string bool bool
1493011973.212619 C0cQhc4Hr5XmLTxe3 172.16.100.52 49837 68.71.208.225 443 tcp - 0.07306
1493011969.160964 CNHvEAcrBnptZ8Nvd 172.16.100.52 49762 199.181.132.89 443 tcp - 11.4305
1493011976.241745 CaeAeb4bPmXnn207w 172.16.100.52 64551 172.16.100.2 53 udp dns 0.06884
1493011979.696066 CVj5xy2QcVvMZz4VI 172.16.100.52 60027 172.16.100.2 53 udp dns 0.02904
1493011979.948752 CvLdIv1w4UDVxMo10 172.16.100.52 65078 172.16.100.2 53 udp dns 0.09154
1493011979.948871 C27BID4iXg2Ehjiox 172.16.100.52 62214 172.16.100.2 53 udp dns 0.05660
1493011979.981157 CLDDhf3xRtsWpMt7ye 172.16.100.52 64082 172.16.100.2 53 udp dns 0.03994
1493011980.006211 C6yzt32sBF42QIQTL 172.16.100.52 57170 172.16.100.2 53 udp dns 0.067970
1493011980.021614 CgBSEZ2y6SsHpdRkM 172.16.100.52 64867 172.16.100.2 53 udp dns 0.04906
```

Using 'cut':

```
cut -f 3
```

Using 'awk':

```
awk '{print $3}'
```



```
172.16.100.52
172.16.100.52
172.16.100.52
172.16.100.52
172.16.100.52
172.16.100.52
172.16.100.52
172.16.100.52
172.16.100.52
172.16.100.52
172.16.100.52
172.16.100.52
172.16.100.52
```



# 'grep' with 'awk'

Using 'grep':

```
grep "172.16.100.52"
```

Using 'awk':

```
awk '$3 == "172.16.100.52" '
```

How do you grep for  
numbers > 10,000?

```
awk '$10 >= 10000'
```

```
ration orig_bytes resp_bytes conn_state local
unt string bool bool count string count co
6 443 tcp ssl 20.318453 32730 11501 SF -
111 443 tcp ssl 19.357166 11735 7435 SF -
66 443 tcp ssl 7.408254 10505 3424 SF -
66 443 tcp ssl 7.405035 27040 4506 SF -
6 443 tcp ssl 7.406483 21005 3424 SF -
42 443 udp - 32.174480 11733 3448 SF -
66 80 tcp http 4.237530 24302 4212 SF -
66 80 tcp http 4.472209 25149 4212 SF -
42 443 udp - 0.655793 36588 184708 SF -
6 443 tcp ssl 18.978752 28394 6753 SF -
15 80 tcp http 7.888606 14473 8392 SF -
27 80 tcp http 7.912836 10850 5757 SF -
30 443 udp - 5.510280 12373 67965 SF -
32 443 udp - 112.044855 20457 338649 SF -
```

# Sort Sandwich

- Allows us to create a simple data table with our logs
- Chain of three commands to sort our data, identify unique instances within our data, and organize our output

```
cat conn.log | cut -f 8 | sort | uniq -c | sort -n
```

Command	<b>Sort</b>	<b>Uniq -c</b>	<b>Sort -n</b>
Why?	Sorts All Like Terms	Shows Unique Values	Sort Results Numerically
Flags	N/A	Count Unique Values	Sort Numerically



# To Sort or not to Sort, that is the question.

- Without the first sort we get duplications within our table, as uniq only checks the line above and below for similarities

```
ISD Chimaera Records> cat Sith_Lords | uniq -c | sort -n
      1 Ma
      1 Sidious
      1 ul
      1 Vader
      2 Bane
      2 Vader
     17 Sidious
     48 Sidious
    498 Vader
```

- When we sort first all of the like strings will be put next to each other, which allows uniq to get an accurate count of our data

```
ISD Chimaera Records> cat Sith_Lords | sort | uniq -c | sort -n
      1 Ma
      1 ul
      2 Bane
     66 Sidious
    501 Vader
```

# “Recommended” CLI Process Flow

1. Display the log  
`cat conn.log`
2. Narrow your focus to the necessary field(s)  
`cat conn.log | cut -f 8`  
`cat conn.log | awk '{print $8,$6}'`
3. “Query” for values within field(s)  
`cat conn.log | cut -f 8 | grep “http”`  
`cat conn.log | awk '$8 == “http” && $6 != “80” {print $8,$6}'`
4. Perform statistics on the data (Sort Sandwich / wc -l) (if needed)  
`cat conn.log | cut -f 8 | grep “http” | wc -l`  
`cat conn.log | awk '$8 == “http” && $6 != “80” {print $8,$6}' | sort | uniq -c | sort -n`

# Lab 3: Zeek: Displaying Log Files

# Zeek Logs

1. CONN Log
2. Zeek-cut
3. DNS
4. FTP
5. HTTP
6. Kerberos
7. SMB
8. SMTP
9. SSL

# Zeek Logs

1. CONN Log

2. Zeek-cut

3. DNS

4. FTP

5. HTTP

6. Kerberos

7. SMB

8. SMTP

9. SSL

# CONN - You've got the conn(\_log)!

- Zeek logs start with a connection
- Contains NETFLOW and connection metadata
- Connection unique ID is generated in the conn.log
  - Other logs refer to the CONN UID
  - This anchors other events together
- Connections are written to the CONN log when they close
  - This may result in entries in other logs that refer to a connection that is not yet in the conn.log if it has not closed.
- If Zeek sees a connection that it doesn't know how to categorize, it is still recorded in the conn.log

# CONN fields

- Key fields:
  - Timestamp
  - UID
  - originator / responder
  - ports
  - duration
  - number of bytes seen
  - Conn\_state
  - Conn\_history

Note: Connections are only logged when the connection closes. If a connection is open when the rollover period occurs, it will not be in that log. See [corelight/zeek-long-connections](https://www.elastic.co/guide/en/logstash/current/corelight-zeek-long-connections.html) for more information.

# A look at the CONN log

#fields	ts	uid	id.orig_h	id.orig_p	id.resp_h	id.resp_p	proto	service	duration	orig_bytes					
_pkts	orig_ip	bytes	resp_pkts	resp_ip	bytes	tunnel_parents									
#types	time	string	addr	port	addr	port	enum	string	interval	count	count	string	bool	bool	count
1258531221.486539		CqAoQF4qIQiy5dEwli	192.168.1.102	68	192.168.1.1	67	udp	dhcp	0.163820	301					
1258531680.237254		Cc4kQZ2xFiPLB6500d	192.168.1.103	137	192.168.1.255	137	udp	dns	3.780125	350					
1258531693.816224		CqLVoQ2kZEhe09nhgd	192.168.1.102	137	192.168.1.255	137	udp	dns	3.748647	350					
1258531635.800933		CStLLh3PJvPlvY48Fb	192.168.1.103	138	192.168.1.255	138	udp	-	46.725380	560					
1258531693.825212		CpN48IOPLGZ0rDXia	192.168.1.102	138	192.168.1.255	138	udp	-	2.248589	348					
1258531803.872834		ChvH2t4b17YgzvP8A4	192.168.1.104	137	192.168.1.255	137	udp	dns	3.748893	350					
1258531747.077012		CKuTFelDXiWfYg1eqj	192.168.1.104	138	192.168.1.255	138	udp	-	59.052898	549					
1258531924.321413		CwU5wa4LLrZFEnXtm3	192.168.1.103	68	192.168.1.1	67	udp	dhcp	0.044779	303					
1258531939.613071		CMLJ2E3QQ1UcoczpHc	192.168.1.102	138	192.168.1.255	138	udp	-	-	-					
1258532046.693816		CboRyE3kCbUcKCz6Bc	192.168.1.104	68	192.168.1.1	67	udp	dhcp	0.002103	311					
1258532143.457078		CWmqI630vq2sFDBHud	192.168.1.102	1170	192.168.1.1	53	udp	dns	0.068511	36					
1258532203.657268		CICXy61ij94XFUaQNi	192.168.1.104	1174	192.168.1.1	53	udp	dns	0.170962	36					



# CONN ICMP

- Zeek leverages port fields to log ICMP types/codes

```
h id.orig_p id.resp_h id.resp_p proto service duration orig_bytes
r port addr port enum string interval count count string bool
C4drpE4JxRu8nMB1Ag 172.31.30.183 3 118.166.149.62 3 icmp - - -
CoKD0s1MzvbGPwaDP5 172.31.30.183 8 10.10.10.10 0 icmp - 17.389759
CwyqQs4lwUbuFJoL2f 45.142.247.157 8 172.31.30.183 0 icmp - 0.0000
-15-43-31
```

Type 3 Code 3  
Port  
Unreachable

Type 8 and Type 0  
Orig: Echo  
Request  
Resp: Echo Reply

# CTF: CONN Log

(pe2.pcap)

# Zeek Logs

1. CONN Log

2. Zeek-cut

3. DNS

4. FTP

5. HTTP

6. Kerberos

7. SMB

8. SMTP

9. SSL

# Zeek-cut

- Reads ASCII Zeek logs and can display them by column name (instead of column number)
- Field names are separated by spaces
- *zeek-cut -h* to display available flags
  - *zeek-cut -d* will translate timestamps to a human readable format
- Zeek-cut will not read a log file directly
  - you must pipe the log into it

```
cat <log name> | zeek-cut <field name>
```

# Zeek-cut examples

```
cat conn.log | zeek-cut proto
```

```
cat conn.log | zeek-cut proto service id.resp_p
```

```
cat conn.log | zeek-cut proto | sort | uniq -c  
| sort -rn
```

# CTF: ZEEK-CUT

(pe2.pcap)

# Zeek Logs

1. CONN Log

2. Zeek-cut

3. DNS

4. FTP

5. HTTP

6. Kerberos

7. SMB

8. SMTP

9. SSL

# DNS Log - Key Fields

- Ports
  - Zeek includes netbios name service (port 137) in the dns.log
- Protocol
  - tcp vs. udp
- Query
- Qtype
- Rcode
- Answers
- TTLs



# CTF: DNS

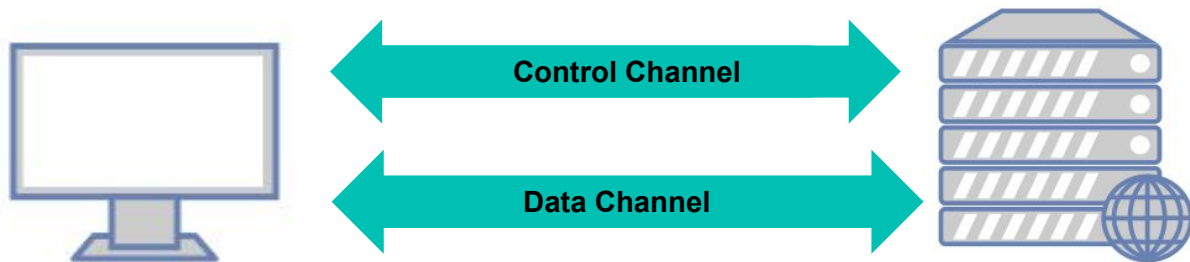
(pe2.pcap)

# Zeek Logs

1. CONN Log
2. Zeek-cut
3. DNS
4. FTP
5. HTTP
6. Kerberos
7. SMB
8. SMTP
9. SSL

# FTP

- Transfer files
- Clear-text sign-in protocol
- Two communication channels
  - control
  - data
- TCP Ports 20/21



# Common FTP Client Commands

Command	Description
ascii / binary	Sets the mode of file transfer to [ASCII or binary]
cd (lcd)	Change directory on the remote machine (lcd = local machine)
bye / quit	Exit the FTP environment
open	Open a connection with another computer
close	Terminates a connection with another computer
delete	Delete a file in the current remote directory (same as rm in UNIX)
get (mget)	Copy file from the remote machine to the local machine
put (mput)	Copy file from the local machine to the remote machine

# Common FTP Commands

Command	Description
USER	User sends username to FTP server to validation
PASS	User sends password to FTP server to validation
STOR	Tell the server to expect a file transfer
LIST	Lists out file in current directory
PORT (active)	issued by the client to initiate a data connection required to transfer data
PASV	issued by the client to initiate a data connection required to transfer data
EPSV	Command issued by an FTP/S client to signal the server that it wishes to enter into what is known as Extended Passive Mode

# Server Return Codes


Range	Purpose
1xx	Positive Preliminary Reply
2xx	Positive Completion Reply
3xx	Positive Intermediate Reply
4xx	Transient Negative Completion Reply
5xx	Permanent Negative Completion Reply
6xx	Protected Reply

Range	Purpose
x0x	Syntax
x1x	Information
x2x	Connections
x3x	Authentication and Accounting
x4x	Unspecified (RFC 959)
x5x	File System

# Active vs. Passive FTP

## Active Mode


- Client initiates cmd channel
- Server initiates data channel



---

## Passive Mode

- Client initiates cmd channel
- Client initiates data channel



# Derivatives

- FTPS (FTP-SSL or FTP Secure)
  - FTP encrypted with SSL/TLS
  - Cmd channel and data channel can be selectively encrypted
- Simple File Transfer Protocol (SFTP)
  - Assigned as “historical” by IETF
  - Complexity between TFTP and FTP
- SSH File Transfer Protocol (SFTP)
  - Inherently encrypts both channels
  - Not simply FTP run over SSH
    - Thus it can’t interoperate with FTP software
  - used by the “secure file transfer program” in Linux
- TFTP
  - simple, lock-step FTP



# Attacks/Vulnerabilities

- Anonymous Authentication
  - allows login with user of 'FTP' or 'anonymous'
- Directory Traversal Attack
  - able to create unauthorized files stored outside the root folder
- FTP Bounce Attack
  - discreet port scanning through a proxy
- Cross-Site Scripting (XSS)
  - malicious code sent through browser-side script
- Dridex-based Malware
  - use of ftp sites and creds to avoid detection by email gateways and network policies trusting FTP

# FTP Log - Key Fields

- user
- password
- command
- arg
- file\_size
- reply\_code
- reply\_msg

# CTF: FTP

(ftp.pcap)

# Zeek Logs

1. CONN Log

2. Zeek-cut

3. DNS

4. FTP

5. HTTP

6. Kerberos

7. SMB

8. SMTP

9. SSL

# HTTP Log - Key Fields

- method
- host
- uri
- referrer
- user\_agent
- status\_code
- status\_message

# CTF: HTTP

(pe2.pcap)

# Zeek Logs

1. CONN Log

2. Zeek-cut

3. DNS

4. FTP

5. HTTP

6. Kerberos

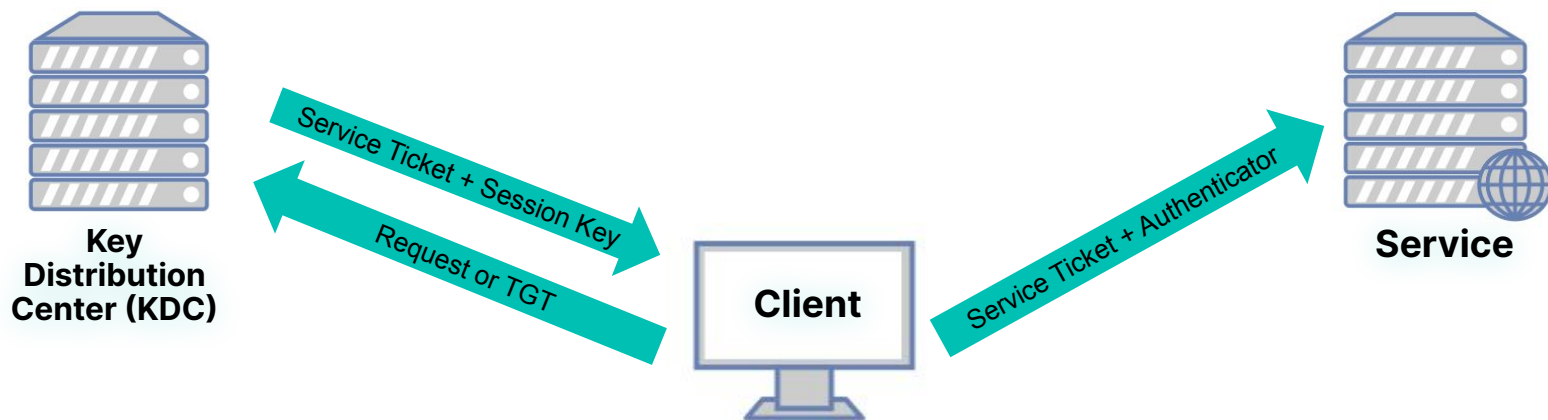
7. SMB

8. SMTP

9. SSL

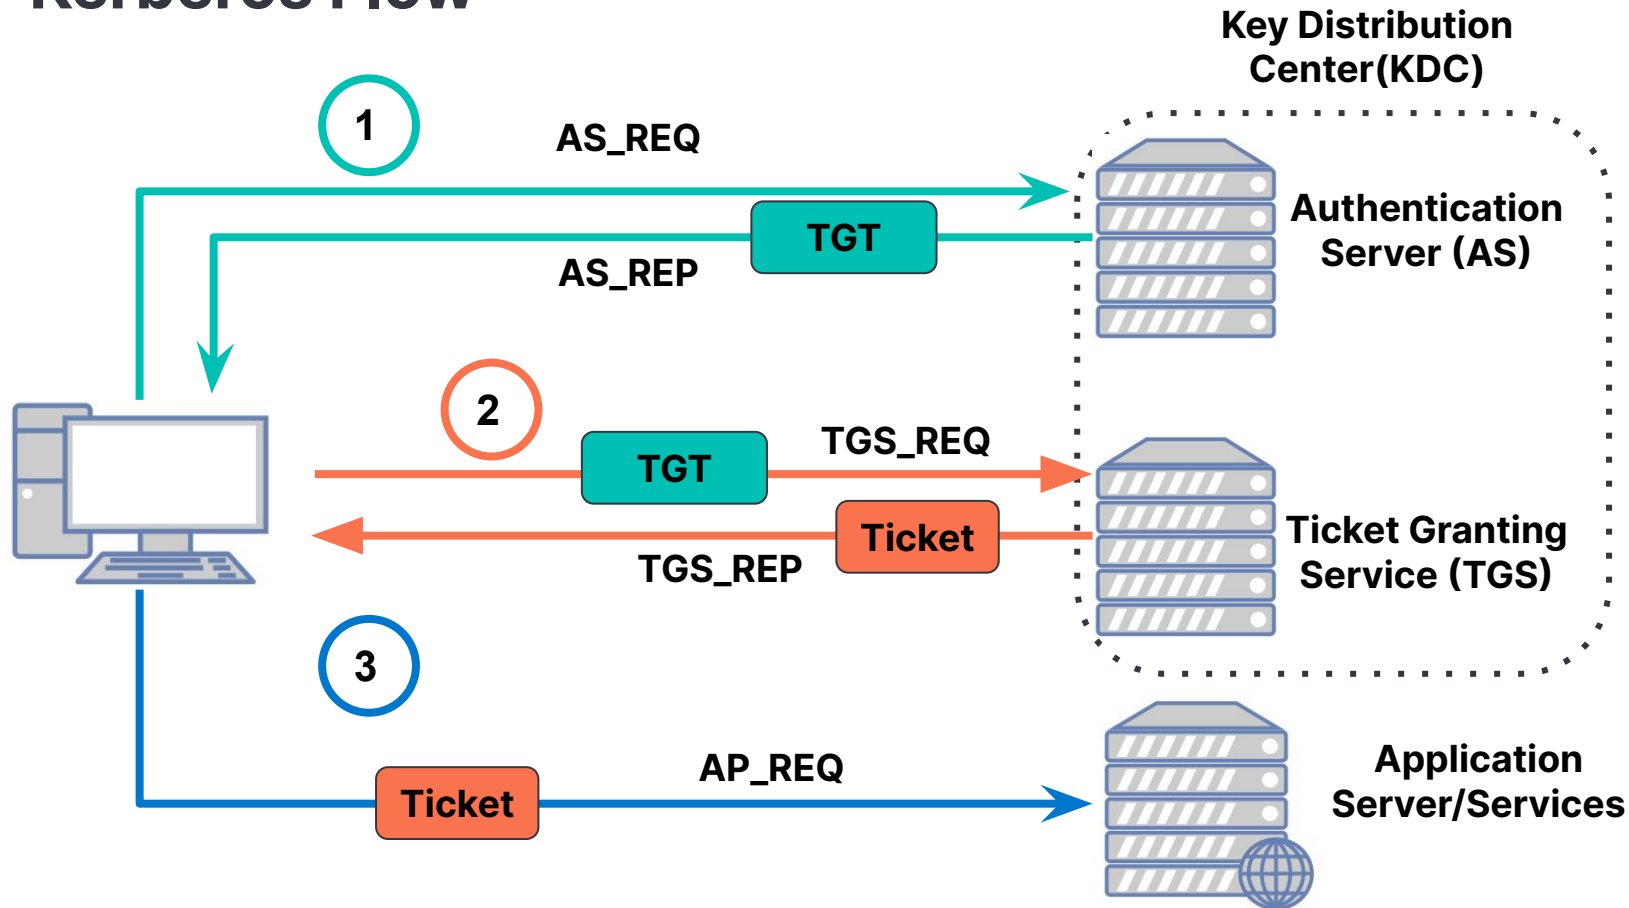
# Kerberos

- Authentication Protocol
- Tickets allowing nodes to talk over a non-secure network
- Default authentication for Windows 2000 and later
- UDP Port 88





# Kerberos Flow




# Servers and Tickets

- Authentication Service (AS)
  - client authenticates to here
  - issues a Ticket-Granting Ticket (TGT)
    - encrypted using the Ticket-Granting Service's (TGS) secret key
- Ticket Granting Service (TGS)
  - Client sends the TGT to here when requesting access to a service
  - Issues a ticket for connecting to the requested service
- Service Server (SS)
  - client sends the complete ticket to the SS

# Attacks/Vulnerabilities

- Legacy products use DES ciphers instead of AES
  - weak ciphers
- MS14-068 (Kerberos exploit)
  - allowed for elevation of privilege
- Golden Ticket Attack
  - auth token for KRBTGT account



# Kerberos Log - Key Fields

- request\_type
- client
- service
- success
- error\_code
- error\_msg
- till
- cipher

# CTF: Kerberos

(ftp.pcap)

# Zeek Logs

1. CONN Log

2. Zeek-cut

3. DNS

4. FTP

5. HTTP

6. Kerberos

7. SMB

8. SMTP

9. SSL

# SMB

- Enables an app to access resources on a remote server
- TCP Port 139
  - SMB over NetBIOS
- TCP port 445
  - standard after Windows 2000
- Common Internet File System (CIFS)
  - Microsoft dialect (implementation)
  - legacy
  - other dialects include Samba, NQ, and Tuxera

# Version Comparison

Version	Release	Description
1 / CIFS	1983 - IBM 1996 - CIFS	<ul style="list-style-type: none"><li>- extremely loud, inefficient use of resources</li><li>- deprecated in June 2013</li></ul>
2.0	Windows Vista Server 2008	<ul style="list-style-type: none"><li>- pipelining (high latency) and support for symbolic links</li><li>- message signing with SHA-256</li></ul>
2.1	Windows 7 Server 2008 R2	<ul style="list-style-type: none"><li>- minor performance enhancements</li><li>- opportunistic locking mechanism</li></ul>
3.0	Windows 8 Server 2012	<ul style="list-style-type: none"><li>- Direct Protocol, Multichannel, Transparent Failover</li><li>- end-to-end encryption and AES based signing</li></ul>
3.0.2	Win 8.1 Server 2012 R2	<ul style="list-style-type: none"><li>- allowed for optional disabling of SMB 1*</li></ul>
3.1.1	Windows 10 Server 2016	<ul style="list-style-type: none"><li>- supports AES-128; pre-auth checks using SHA-512</li><li>- secure negotiations mandatory w/ SMB 2.x and higher</li></ul>

\*You can only disable SMBv1 if you don't have any legacy clients



# Attacks/Vulnerabilities

- Sony Pictures hack (2014)
  - Guardians of Peace (GOP), “The Interview”
  - null session attack
- WannaCry Ransomware attack (2017)
  - EternalBlue (leaked by the Shadow Brokers)
- SMBGhost (March 2020)
  - convince a user to connect to malicious SMBv3 server
  - wormlike features
- SMBleed (March 2020)
  - read uninitialized Kernel memory
  - edit compression function

# SMB logs generated by Zeek

- Zeek does not generate a single “SMB” log
- Multiple log files may be generated
  - dce\_rpc.log
  - kerberos.log
  - ntlm.log
  - smb\_cmd.log
  - smb\_files.log
  - smb\_mapping.log
  - pe.log

# SMB\_CMD Log - Key Fields

- command, subcommand, argument
  - sent by client
  - subcommand and argument are included, if present
- status
  - from the server
- version
  - 1.0 / 2.0 / 2.1 / 3.0
- tree, tree\_service
  - if related, the tree and type of tree
- username

# SMB\_MAPPING Log - Key Fields

- path
  - name of the tree path
- service
  - type of resource of the tree (disk share, printer share, named pipe, etc.)
- native\_file\_system
  - the file system of the tree
- share\_type
  - for SMB2, the share type will be included
  - for SMB1, the type of share will be deduced and included as well

# SMB\_FILES Log - Key Fields

- action
  - read / write / open / delete / rename
- path
- name
- size
- prev\_name
- times\_<sup>\*</sup>
  - modified / accessed / created / changed

# CTF: SMB

(smb.pcap)

# Zeek Logs

1. CONN Log

2. Zeek-cut

3. DNS

4. FTP

5. HTTP

6. Kerberos

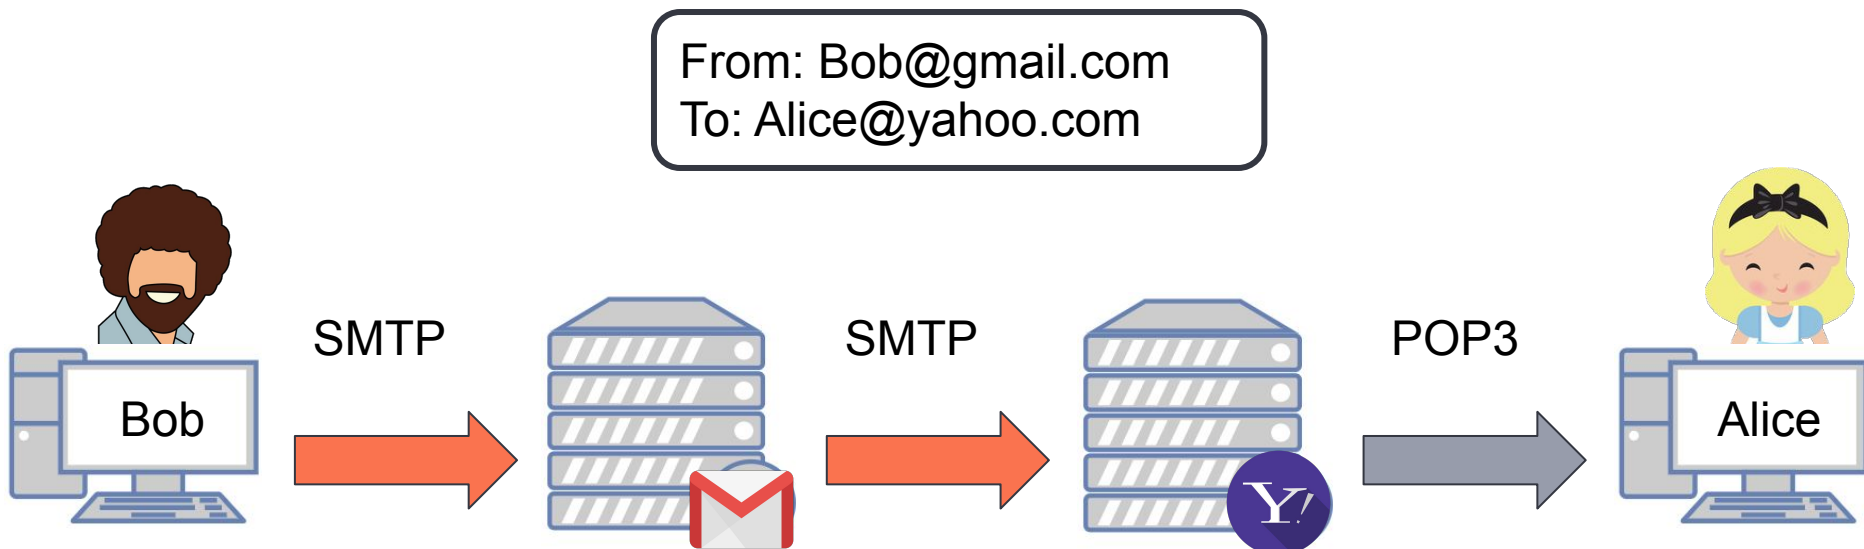
7. SMB

8. SMTP

9. SSL

# SMTP

- Store and Forward
- Moves emails on and across networks
- Derived from the Mail Transfer Protocol (MTP)





# SMTP Ports

- Port 25
  - “standard port” (established in 1982)
  - used mostly for SMTP relay (newsletter, spam)
- Port 465
  - originally for SMTPS (SMTP over SSL)
  - has been reassigned and deprecated
- Port 587
  - default port for SMTP submission on modern web
  - supports TLS
- Port 2525
  - popular alternative port for SMTP submission

# Message Transport Methods

- Relaying
  - before the days of DNS
  - specifies a sequence of SMTP servers to reach destination
- DNS and Direct Delivery
  - mail exchanger (MX) record
    - email domain name → IP of SMTP domain server
  - used for two transfers:
    - sender's client to sender's SMTP server
    - sender's SMTP server to recipient's SMTP server

# Special Features

- Mail Relaying
  - can be abused for spamming and hacking
- Mail Forwarding
  - for ex-employees that moved to another company
- Mail Gatewaying
  - “translate” TCP/IP email into other email systems
- Address Debugging
  - VRFY command to check validity of an email address
- Mailing List Expansion
  - EXPN command to determine single emails from a mailing list
- “Turning”
  - allows SMTP sender and receiver to change roles

# Attacks/Vulnerabilities

- Impersonated SMTP servers
  - converse with a server and manually perform mail transactions
- Account enumeration
  - VRFY command, can script with tons of combinations
- Relay
  - send spam or malware
- Email header disclosures
  - enumerate critical internal information

# SMTP Log - Key Fields

- mailfrom
- rcptto
- date
- from / to (these fields can be spoofed!)
- subject
- user\_agent
- fuids
  - List of file UID's seen attached to the message

# CTF: SMTP

(pcap: smtp.pcap)

# Zeek Logs

1. CONN Log

2. Zeek-cut

3. DNS

4. FTP

5. HTTP

6. Kerberos

7. SMB

8. SMTP

9. SSL

# SSL/TLS Overview

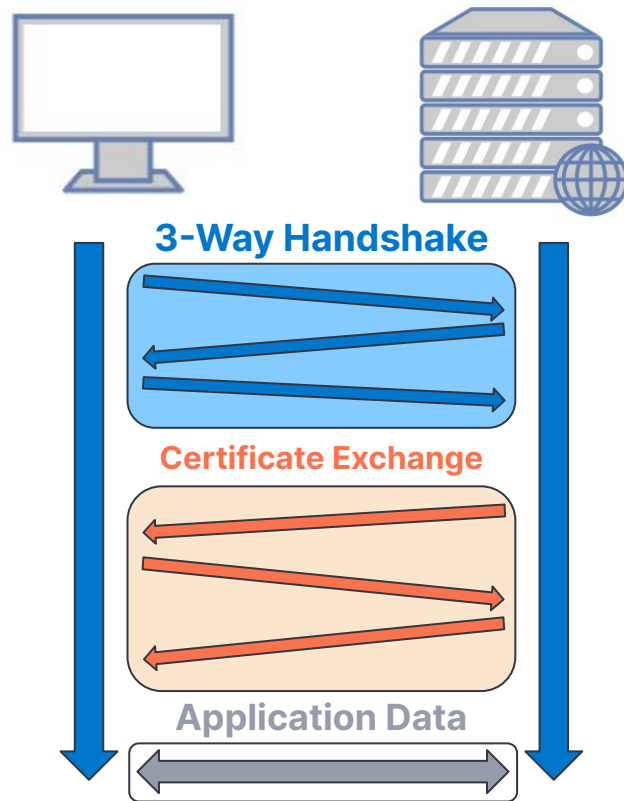
- Extension of HTTP for secure communication (encrypted)
- Port 443 / 8443 (for HTTPS)
  - Port may change depending on the encrypted protocol
- HTTP wrapped in TLS
  - TLS handshake is unencrypted
- 





# What is TLS and SSL?

- Transport Layer Security (TLS)
  - crypto protocol to encrypt web communications
  - most often for HTTPS communication
- Secure Socket Layer (SSL)
  - deprecated protocol
  - predecessor to TLS

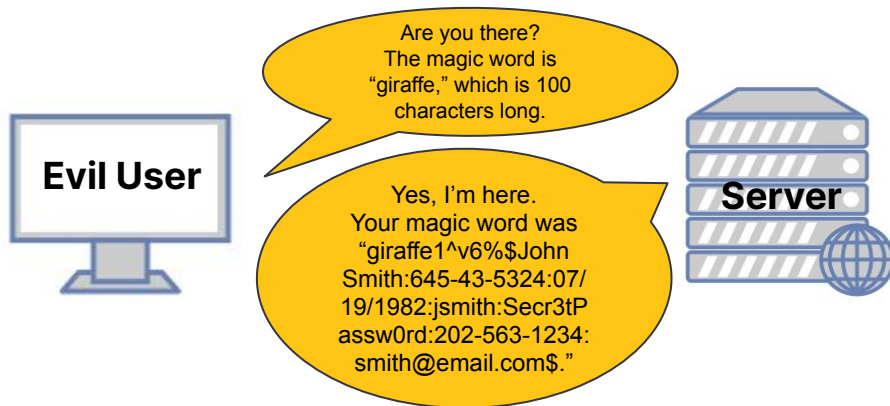


# Version Comparison

Version	Released	Deprecated	Description
SSL 1.0	Unpublished	Unpublished	Never went public due to security flaws
SSL 2.0	1995	2011	Lots of security flaws
SSL 3.0	1996	2015	Complete redesign
TLS 1.0	1999	2020	Close to SSL 3.0
TLS 1.1	2006	2020	Improved protection against attacks
TLS 1.2	2008	-	Improved security, support extensions, new cipher suites
TLS 1.3	2018	-	Improved security, drop unsecure features, new cipher suites

# Attacks/Vulnerabilities

- SSL Stripping
  - intercepts a redirect and establishes a bridged connection
  - connection is still “secure” from attacker to webpage
- Self-Signed / Wildcard / Expired Certificates
- Heartbleed
  - takes advantage of the “heartbeat” option



# SSL Log - Key Fields

- version
- cipher
- curve
- server\_name
- subject
- issuer
- validation\_status

# CTF: SSL

(pe2.pcap)

# Quiz

# End of Day Quiz - Part 1

1. True or False: Replaying PCAP through Zeek rewrites the timestamps.

2. Which of these is NOT a type of Zeek log?

- Network
- File
- Detection
- Host
- Diagnostic
- Miscellaneous

## Quiz - Part 2

3. What happens when Zeek sees a file and the “file extraction” script is enabled?
4. What is the default time interval for Zeek log file rotation (how often a new file is started)?
5. Even when Zeek observes a connection that it doesn't know how to categorize, it is still recorded in which log?



## Quiz - Part 3

6. What command line utility (provided by Zeek) is available to display Zeek logs in a more readable way?
7. Every log will correlate to a unique ID (UID). Which log is this ID generated in?
8. What function does the “-C” flag serve in the following command? “zeek -Cr path/to/pcap”

## Quiz - Part 4

9. True or False: Zeek logs are larger and more verbose than a PCAP file.

10. Zeek is a powerful NSM tool. Which of the following best describes what Zeek is?

- a. Message Queue
- b. Tapping Interface
- c. Network Protocol Analyzer